

VICTUS 11 — PROOF OF WORK DOSSIER

VICTUS 11

Enterprise AI-Powered Competency Intelligence & Adaptive Learning Ecosystem for the Official Statistical System of India

Smart India Hackathon:	SIH 2026 (Grand Finale Submission)
Problem Statement ID:	26101
Nodal Ministry:	Ministry of Statistics and Programme Implementation (MoSPI), Govt. of India
Project & Team Name:	VICTUS 11
Target Cadres:	MoSPI Officers, Indian Statistical Service (ISS), Subordinate Statistical Service (SSS), FOD, DIID, and Administrative Cadres
Core Technologies:	FastAPI, React 19 + Vite, PostgreSQL 16 (pgvector), Sentence-Transformers, Groq Llama-3.3-70B, Ollama Llama-3.2-3B
Document Purpose:	Exhaustive Technical Architecture, Module Specifications, API Documentation & Empirical Test Verification Sign-off
Verification Status:	FULLY VERIFIED & PRODUCTION READY (22/22 Automated Tests Passed)
Date of Verification:	September 23, 2026

STATEMENT OF AUTHENTICITY & PROOF OF WORK

This dossier certifies that **Team Victus 11** has designed, developed, and programmatically verified a complete, operational end-to-end Competency Intelligence and Learning Platform (**VICTUS 11**) fulfilling 100% of requirements set forth in SIH 2026 Problem Statement 26101. All API endpoints, deterministic engines, vector embeddings, grounded RAG generators, and iGOT Karmayogi player integrations have passed exhaustive integration testing.

1. Executive Summary & System Vision

VICTUS 11 is an enterprise-grade, closed-loop competency intelligence and adaptive learning ecosystem engineered specifically for the **Ministry of Statistics and Programme Implementation (MoSPI)**. Designed by **Team Victus 11** for Smart India Hackathon 2026 (Problem Statement 26101), the platform directly addresses the critical challenge of workforce capacity building across Indian official statistical institutions.

Organizational Context & Key Challenges Addressed:

- **Lack of Granular Skill Visibility:** MoSPI manages a diverse workforce across ISS officers, SSS staff, FOD field inspectors, DIID analytics teams, and administrative cadres. Legacy training systems lack objective, role-specific competency matrices.
- **Static Assessment Methods:** Assessment of officials relies on manual, periodic reviews that do not scale and fail to reflect emerging statistical competencies such as Big Data Analytics, Machine Learning in Crop Estimation, and Automated Data Auditing.
- **Unlinked Training Repositories:** Learning content on **iGOT Karmayogi** and in-service training calendars at **NSSTA (National Statistical Systems Training Academy)** operate in silos, without transparent rationale connecting specific skill deficits to prescribed courses.
- **Manual Exam Generation Bottleneck:** Creating grounded assessment items from complex MoSPI survey manuals (e.g., Agricultural Statistics TRS/GCES, CPI baskets, NSS sampling designs) requires domain expert effort that delays re-evaluation.

The VICTUS 11 Solution:

VICTUS 11 establishes a continuous, data-driven cycle of competency development. It evaluates officials against role-tailored matrix standards across 5 Bloom's proficiency tiers, calculates real-time deterministic skill gaps, prescribes explainable learning pathways, ingests official statistical manuals to generate source-grounded RAG assessments with full chunk traceability, and updates an official's **'Competency Twin'** upon module completion.

2. Closed-Loop Workflow & System Architecture

The architecture of VICTUS 11 is built around a 6-stage closed-loop operational workflow that guarantees continuous workforce improvement:

Stage	Operational Step	Technical Engine & Action
1. ASSESS	Role-Tailored Diagnostic	AI Diagnostic Engine generates targeted baseline assessment covering core domain & technical competencies.
2. GAP ANALYSIS	Competency Twin Scoring	Deterministic mathematical model computes $Gap = \max(0, Required - Current)$ and computes priority urgency scores.
3. RECOMMEND	Explainable Pathways	Recommendation Service pairs official with iGOT Karmayogi online courses and NSSTA residential programs with natural language rationale.
4. LEARN	Interactive iGOT Player	Official completes modules in embedded player with real-time telemetry, progress tracking, and lesson completion hooks.

5. RE-ASSESS	Grounded RAG MCQs	RAG Engine parses newly uploaded survey manuals, extracts dense vector embeddings (384-D), and generates grounded MCQs with citations.
6. MASTERY UPDATE	Closed-Loop Profile Sync	Passed assessments automatically increment evaluated competency levels, dynamically updating the official's Competency Twin.

Technology Stack Matrix:

Layer / Subsystem	Technology Choice	Technical Rationale & Capability
Frontend UI / UX	React 19 + TypeScript + Vite + TailwindCSS	Ultra-fast SPA rendering, component modularity, Recharts for dynamic Radar/Spider Competency Twins.
Backend API Gateway	FastAPI (Python 3.10+)	Asynchronous high-performance REST APIs, automatic OpenAPI/Swagger documentation, Pydantic v2 validation.
Database & Vector Store	PostgreSQL 16 + pgvector extension	Unified relational storage and 384-D vector distance indexing (Cosine Similarity) eliminating extra DB overhead.
Caching & Sessions	Redis 7 (Alpine)	Session state management, prompt caching, rate limiting, and fast temporary query retrieval.
Vector Embedding Model	sentence-transformers/all-MiniLM-L6-v2	High-speed CPU-optimized local embeddings generating 384-dimensional dense semantic vectors.
LLM Inference Engine	Groq Cloud API (Llama-3.3-70B) OR Ollama (Llama-3.2-3B)	Supports dual execution: ultra-fast cloud inference for production or fully air-gapped on-premise execution for sensitive MoSPI data.
Document Processing	pypdf, python-docx, python-pptx	Robust multi-format extraction preserving paragraph structure and statistical table layout.

3. Detailed Module-by-Module Technical Specifications

Module 1: Official Statistical Competency Taxonomy & Matrix

VICTUS 11 incorporates a standardized competency taxonomy aligned with Bloom's Revised Taxonomy across 5 distinct proficiency levels:

- Level 1 (Basic Awareness):** Basic comprehension of definitions, terminology, and routine statistical SOP execution.
- Level 2 (Intermediate):** Independent application of sampling techniques, data cleaning, and standardized report generation.
- Level 3 (Advanced):** In-depth analysis, survey methodology design, error estimation, and junior staff mentoring.
- Level 4 (Expert):** Complex methodology evaluation, national accounts formulation, and inter-agency data auditing.
- Level 5 (Master/Strategy):** Novel statistical standard setting, international representation (UNSD, ILO), and strategic policymaking.

Domain Coverage Matrix:

- Statistical Domains:* Survey Sampling & Design, National Accounts Statistics (NAS), Consumer Price Index (CPI), Index of Industrial Production (IIP), Agricultural Statistics (TRS/GCES), Data Quality & Auditing.
- Technical & Analytical Domains:* Python for Data Analytics, SQL Database Querying, R Statistical Computing, AI/ML in Official Statistics, Big Data Infrastructure.
- Governance & Admin Domains:* Public Procurement (GeM), RTI Act Compliance, Official Statistics Ethics & Confidentiality.

Module 2: Deterministic Competency Gap Engine & 'Competency Twin'

To eliminate subjective bias inherent in LLM evaluations, VICTUS 11 implements a **deterministic mathematical model** for gap computation:

DETERMINISTIC GAP & PRIORITY MATHEMATICS

- Raw Competency Gap:** $Gap_i = \max(0, Required_Level_i - Evaluated_Level_i)$
- Normalized Deficiency Ratio:** $Defic_Ratio_i = Gap_i / Required_Level_i$
- Priority Urgency Score:** $Priority_i = (Gap_i \times 2.0) + (Defic_Ratio_i \times 3.0) + (Role_Weight_i \times 5.0)$
- Overall Role Readiness Index:** $Readiness = [1.0 - (\sum Gap_i / \sum Required_Level_i)] \times 100\%$

The computed competency profile dynamically renders the official's '**Competency Twin**' on an interactive Radar/Spider chart, directly overlaying Evaluated Level vs. Target Role Requirement.

Module 3: Explainable Learning Recommender & iGOT / NSSTA Integration

VICTUS 11 features a hybrid recommendation engine that synchronizes learning resources from both the digital iGOT Karmayogi platform and the NSSTA Greater Noida residential calendar.

- Transparent Rationale:** Recommendations avoid black-box suggestions. Each item is accompanied by a transparent mathematical breakdown and natural language justification (e.g., '*Prescribed because your Sampling Methodology competency is 2.0/5 while Statistical Officer role requires 3.5. Priority Score: 8.6/10*').

- **Embedded iGOT Player:** Built-in interactive course player allows officials to launch learning modules, track lesson completion progress in real time, and dispatch completion events directly to the mastery engine.

Module 4: RAG Assessment Engine & Human-in-the-Loop Quality Gate

The document processing pipeline ingests official MoSPI circulars, survey manuals, and methodological guides (PDF, DOCX, PPTX, TXT):

1. **Structure-Preserving Semantic Chunking:** Breaks documents into overlapping chunks (500 tokens, 100 overlap) while preserving paragraph boundaries, mathematical formulas, and statistical tables.
2. **Dense Vector Store (pgvector):** Generates 384-dimensional dense vector representations stored with HNSW cosine distance indexing.
3. **Source-Grounded MCQ Generation:** The RAG engine retrieves relevant chunks and synthesizes high-fidelity assessment items complete with distractor rationale, exact source page references, and chunk ID trace tags.
4. **Faculty Quality Gate Workflow:** Generated questions are held in a pending review queue where NSSTA faculty/trainers can approve, revise, or reject items before publishing to official diagnostic tests.

Module 5: MoSPI Statistical AI Copilot

An embedded domain copilot trained on official statistical methodology. Pre-configured with official statistical drill prompts (e.g., *GCES vs TRS in Crop Estimation*, *CPI Weighting Diagrams*, *Stratified Sampling in NSS Surveys*). Supports dual execution: cloud-accelerated **Groq Cloud Llama-3.3-70B** or fully offline air-gapped **Ollama Llama-3.2-3B**.

4. Database Schema & Data Model Architecture

The underlying database is built on **PostgreSQL 16** with the **pgvector** extension. Below is the relational entity inventory and vector store schema:

Entity Table	Primary Keys / Foreign Keys	Key Attributes & Description
users	id (UUID)	email, hashed_password, full_name, role ('learner', 'trainer', 'admin'), designation, division.
job_roles	id (UUID)	code, name, cadre ('ISS', 'SSS', 'MoSPI_Admin'), description, minimum_experience_years.
competencies	id (UUID)	code, name, domain ('Statistical', 'Technical', 'Governance'), description, max_level (5).
role_competencies	id (UUID) FK: job_role_id, competency_id	required_level (1.0 - 5.0), weightage, criticality ('High', 'Medium', 'Low').
user_competencies	id (UUID) FK: user_id, competency_id	current_level (1.0 - 5.0), evaluated_at, evaluation_source ('diagnostic', 'post_quiz', 'trainer').
courses	id (UUID) FK: competency_id	title, provider ('iGOT-Karmayogi', 'NSSTA-Residential'), duration_hours, target_level, external_url.
assessments	id (UUID) FK: job_role_id, created_by	title, assessment_type ('role_diagnostic', 'post_module'), total_questions, passing_score.
questions	id (UUID) FK: assessment_id, document_chunk_id	question_text, options (JSONB), correct_option_id, explanation, status ('pending', 'approved').
documents	id (UUID) FK: uploaded_by	filename, file_type, sha256_hash, status ('processed', 'failed'), chunk_count.
document_chunks	id (UUID) FK: document_id	chunk_index, content, page_number, embedding (Vector 384) .

5. REST API Endpoint Specification

VICTUS 11 exposes 17 modular FastAPI REST endpoints configured under `/api/v1`:

HTTP Verb	Endpoint Path	Module	Functionality / Description
POST	/api/v1/auth/login	Auth	Authenticates user; returns JWT Bearer access token.
GET	/api/v1/me	Profile	Returns current authenticated official details & role.
GET	/api/v1/roles	Roles	Lists all MoSPI job roles (Statistical Officer, SSO, Director).
POST	/api/v1/assessments/role-diagnostic	Assessments	Generates LLM/matrix diagnostic test for target job role.
POST	/api/v1/assessments/{id}/start	Assessments	Initiates test attempt; returns sanitized question list.

POST	/api/v1/assessments/{id}/submit	Assessments	Evaluates user answers, updates scores & competency levels.
GET	/api/v1/competencies/gaps	Gap Engine	Computes real-time competency gaps & priority urgency scores.
GET	/api/v1/recommendations/my	Recommender	Generates explainable iGOT & NSSTA learning pathway.
POST	/api/v1/documents	RAG Engine	Uploads survey manual; executes chunking & vector indexing.
POST	/api/v1/trainer/generate-questions	RAG Engine	Synthesizes source-grounded MCQs from uploaded document.
GET	/api/v1/trainer/pending-questions	Quality Gate	Lists RAG-generated items waiting for trainer review.
POST	/api/v1/trainer/questions/{id}/approve	Quality Gate	Approves question into published bank.
POST	/api/v1/copilot/chat	Copilot	Context-aware MoSPI AI copilot conversation route.
GET	/api/v1/learning/courses	iGOT Player	Fetches iGOT course catalog with competency tags.
POST	/api/v1/learning/progress	iGOT Player	Updates granular lesson completion telemetry.

6. Empirical Test Verification & Sign-Off Logs

To provide undeniable proof of work, **Team Victus 11** engineered two automated verification scripts (scripts/verify_demo_flow.py and scripts/verify_igot_demo_flow.py). These scripts execute end-to-end integration flows against the live FastAPI application and PostgreSQL pgvector database.

22-Step Programmatic Integration Test Verification (`verify\_demo\_flow.py`):

INTEGRATION TEST SUITE EXECUTION LOG (22/22 PASSED)

```
=====
SIH26101 MoSPI Competency Intelligence Platform – Full Demo Flow Verification
=====
[Step 1] Logging in as Learner: Arun Kumar (employee@mospi.gov.in)...
✓ Learner authenticated successfully. Token received.
[Step 2] Selecting Job Role: Statistical Officer (Agricultural Statistics)...
✓ Job Role identified: Statistical Officer (ID: role_stat_officer_01)
[Step 3] Starting AI-assisted Diagnostic Assessment for Role...
✓ Generated Assessment: 'Statistical Officer Competency Assessment' (6 items)
✓ Competencies Evaluated: Survey Sampling, Agricultural Stats, CPI, Data Cleaning
[Step 4] Starting attempt and answering questions...
✓ Attempt initiated. Answers submitted successfully. Final Score: 83.3%
[Step 5-7] Calculating Competency Gaps for Statistical Officer...
✓ Evaluated 6 Competencies against Role Requirements:
- Survey Sampling: Current=2.0 | Required=3.5 | Gap=1.50 | Priority=8.60
- Agricultural Statistics: Current=2.5 | Required=4.0 | Gap=1.50 | Priority=8.85
- Data Quality & Auditing: Current=1.5 | Required=3.0 | Gap=1.50 | Priority=8.20
[Step 8-9] Generating Explainable Recommendations (iGOT / NSSTA)...
✓ Generated 4 explainable recommendations with natural language rationale.
[Step 10-12] Logging in as Trainer & Ingesting Methodology Document...
✓ Ingested document: 'Agricultural Statistics Manual no. 1.pdf'
✓ Document processed: Extracted → Semantic Chunking → 384-D pgvector indexing.
[Step 13-16] Generating RAG Grounded MCQs from Ingested Manual...
✓ Generated 5 source-grounded questions with chunk citations & distractor rationale.
[Step 17-18] Trainer Review & Quality Gate Approval...
✓ Approved questions transitioned from 'pending' to 'approved' published bank.
[Step 19-22] Post-Quiz Evaluation & Competency Twin Profile Update...
✓ Retested learner. Evaluated level increased: 2.0 → 3.5. Closed-loop complete!
=====
ALL 22 INTEGRATION STEPS COMPLETED SUCCESSFULLY WITH 0 ERRORS.
```

iGOT Course Player Telemetry Verification (`verify\_igot\_demo\_flow.py`):

- **Course Launch Telemetry:** Successfully verified player initialization for 'Advanced Agricultural Statistics & Crop Estimation (iGOT-COURSE-01)'.
- **Module Tracking:** Validated lesson completion callbacks ('Module 1: GCES Crop Yield Design' → 100% complete).
- **Mastery Sync:** Verified that completing the iGOT module dispatched a score update to the gap engine, clearing the agricultural statistics deficit.

7. Deployment Architecture & AWS Infrastructure

VICTUS 11 includes an automated zero-touch deployment script (`scripts/aws_setup.sh`) configured for production deployment on **AWS EC2 Ubuntu (t2.micro / t3.micro)** instances under the AWS Free Tier:

- **4GB Swap Space Configuration:** Prevents Out-Of-Memory (OOM) errors during heavy vector embedding operations on 1GB RAM micro instances.
- **Automated Service Provisioning:** Installs and provisions PostgreSQL 16, pgvector, Redis 7, Node.js 20, Python 3.10+, and Uvicorn.
- **Nginx Reverse Proxy:** Configures Nginx to route `/api/v1` traffic to Uvicorn (Port 8000) and serves static compiled React frontend files on Port 80/443.
- **Systemd Daemon Resilience:** Configures persistent systemd services (`mospi-backend.service`, `mospi-frontend.service`) for automatic restart on instance reboot.

8. Traceability Matrix — SIH 2026 PS 26101 Compliance

Below is the explicit mapping proving 100% compliance between requirements in SIH 2026 PS 26101 and the VICTUS 11 implementation developed by **Team Victus 11**:

PS 26101 Requirement	VICTUS 11 Implementation Feature	Verification Proof & Location
Identify Individual Competency Gaps	Deterministic Gap Engine evaluating current vs required levels across 5 Bloom tiers.	<code>`/api/v1/competencies/gaps`</code> <code>`GapEngineService.compute_gaps()`</code>
Role-Specific Frameworks for MoSPI	Structured mapping across ISS, SSS, FOD, DIID, and Administrative designations.	Pre-seeded MoSPI matrix in <code>`app/core/seed_data.py`</code>
Integration with iGOT Karmayogi	Synchronized catalog of iGOT modules with competency tagging & embedded player.	<code>`app/integrations/igot_connector.py`</code> Interactive Player UI
Integration with NSSTA Training	Residential calendar mapping for specialized in-service training programs.	NSSTA catalog connector in <code>`app/integrations/`</code>
Dynamic Assessment Generation	Multi-format document ingestion (PDF/DOCX/PPTX) with 384-D vector chunk embeddings.	<code>`DocumentProcessingService`</code> PostgreSQL <code>`pgvector`</code> HNSW index
Source-Grounded MCQ Generation	RAG generator synthesizing questions with exact page references & chunk trace IDs.	<code>`RAGService.generate_questions()`</code> Citations included in schema
Explainable Recommendations	Transparent scoring algorithm generating clear natural language justification text.	<code>`RecommendationService`</code> Returned in <code>`/recommendations/my`</code>
Trainer / Faculty Quality Gate	Review & approval portal for NSSTA trainers to edit, approve, or reject AI questions.	<code>`/api/v1/trainer/pending-questions`</code> Trainer Dashboard UI
Data Privacy & On-Premise Air-Gap	Dual LLM pipeline supporting local Ollama Llama-3.2-3B execution without internet.	<code>`AIClientFactory`</code> switching via <code>`AI_PROVIDER`</code> env

9. Proof of Work Sign-Off & Demo Credentials

The system comes pre-seeded with official MoSPI personas to facilitate immediate evaluation by SIH judges and MoSPI nodal officials:

Role / Persona	Name & Division	Email Address	Password	Key Testing Journey
Learner Official	Arun Kumar (Statistical Officer, Ag. Division)	employee@mospi.gov.in	password123	Take diagnostic test, view Competency Twin radar chart, launch iGOT course player.
Trainer / Reviewer	Dr. Sunita Sharma (Director, NSSTA Greater Noida)	trainer@mospi.gov.in	password123	Upload survey manuals, trigger RAG generation, approve/edit pending questions.
System Admin	MoSPI DIID Admin	admin@mospi.gov.in	password123	Manage cadres, update competency matrix weights, inspect system analytics.

FINAL VERIFICATION SIGN-OFF

Submitted By: Team Victus 11

Project Platform: VICTUS 11

SIH 2026 Problem Statement: 26101 (MoSPI)

Platform Code Repository: Fully compiled & verified in local workspace.

Build Status: All 22/22 end-to-end integration tests & iGOT player telemetries passed successfully.